A Neuro-Symbolic Benchmark Suite for Concept Quality and Reasoning Shortcuts

S. Bortolotti[1] **E. Marconato**[1,2] T. Carraro[4, 5] P. Morettin[1] E. v. Krieken[3] A. Vergari[3] S. Teso[1] A. Passerini[1]

[1]University of Trento [2]University of Pisa [3]University of Edinburgh [4]Fondazione Bruno Kessler [5]University of Padova

PAPER CODE

REASONING SHORTCUTS

Goal: Study supervised models that classify samples *correctly* but for the *wrong concepts*.

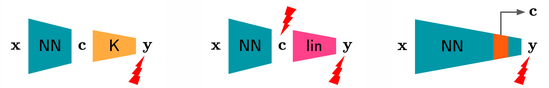

(a) NeSy (DPL & LTN) (b) CBM (c) NN & CLIP

Reasoning Shorcuts [1]: NeSy predictors [2], Concept-based Models [3] and VLMs like CLIP [4] solve Learning & Reasoning tasks by exploiting semantically misleading concepts.

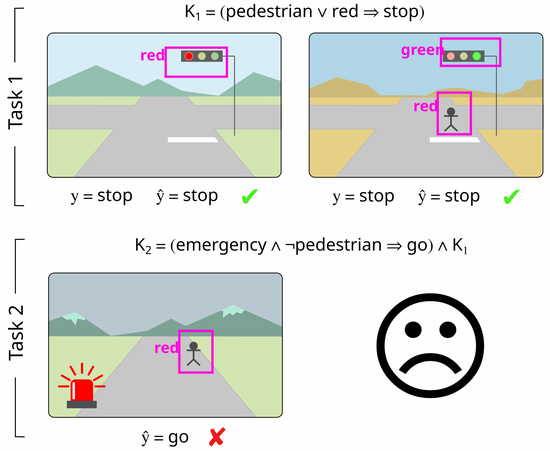

L&R TASKS

	Task	Data: Gen	Data: OOD	Data: Cont	Properties: Cplx x	Properties: Cplx K	Properties: Amb K
	MNMath (new)	✓	✓	✓	✗	✓	✗
+/×	MNAdd-Half	✗	✓✓	✗	✗	✗	–
	MNAdd-EvenOdd	✗	✓✓	✓✓	✗	✗	–
	MNLogic (new)	✓	✓	✓	✗	✓	✗
∧/∨	Kand-Logic	✓	✓	✓	✗	✓	✓
	CLE4EVR	✓	✓	✓✓	✓	✗	✓
⚠	BDD-OIA	✗	✗	✗	✓	✓	✓
	SDD-OIA (new)	✓	✓✓	✓	✓	✓	✓

FEATURES

① **Challenging**: the # of RSs can be chosen **a priori** and counted using `countrss`, allows to control task difficulty.

② **Configurable**: data sets & generators can be easily configured with `YAML/JSON` files.

③ **Intuitive**: straightforward to use:

```
from rsbench import MNLOGIC

dataset = MNLOGIC(args)
train(model, dataset)
test(model, dataset)
```

NOTEBOOK

EXAMPLES

Task	Example	Shortcut	OOD Pred.
SDD-OIA		STOP	GO
BDD-OIA		STOP	GO

Knowledge K = *the traffic laws.*

Task	Example	Shortcut	OOD Pred.
MNMath	$2 \cdot \blacksquare + \blacksquare = 6$, $\blacksquare + \blacksquare = 7$	$\blacksquare \to 2$, $\blacksquare \to 4$, $\blacksquare \to 3$	$\blacksquare + \blacksquare = 5$

Knowledge K = *equations must hold.*

Task	Example	Shortcut	OOD Pred.
MNLogic [1]	$\blacksquare \oplus \blacksquare = 1$	$\blacksquare \to 1$, $\blacksquare \to 0$	$\blacksquare \wedge \blacksquare = 1$
Kand-Logic [2]	$= 1$	□ → red, △ → yel, ○ → blu	$= 0$
CLE4EVR [3]	$= 0$, $= 1$	■ → ■, ■ → ■, ○ → ■	$= 1$

1: Knowledge K = *formula must hold.*
2: Knowledge K = *pattern must hold.*
3: Knowledge K = *same color and shape?*

ASSESSING RS

① **Task-level**: `countrss` counts the # of potential RSs in any L&R task!

$$\mathsf{K} : Y \leftrightarrow (C_1 \wedge C_2 \wedge C_3) \qquad \mathcal{D} : \{(0,1,0),(1,0,0)\}$$

`countrss`

$$\#RSs$$

Example: *with* 3 *concepts and an exhaustive training set,* `MNLogic` *has* 6 *RSs if* K *is a conjunction and* 24 *if* K *is a XOR. This grows* ***exponentially*** *with the # of concepts!*

.

② **Model-level**: `rsbench` tasks induce RSs in all models!

Table 1. (L) DPL and (R) NN concept confusion matrix on `MNAdd-EvenOdd`

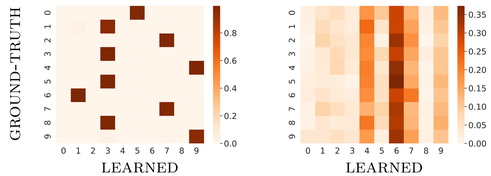

Quantitatively: Concept F1, accuracy and *collapse*

REFERENCES

[1] Marconato *et al.*, Analysis and Mitigation of RSs, NeurIPS (2023)
[2] Manhaeve *et al.*, DeepProblog, NeurIPS (2018)
[3] Pang Wei Koh *et al.*, Concept bottleneck models, ICML (2020)
[4] Alec Radford *et al.*, CLIP, ICML (2021)